

A comparison of Tensorflow and PyTorch from the perspective of a beginner

Kevin M. Loew


July 23, 2019

Outline

- ◎ Background
 - What are TensorFlow and pyTorch?
- ◎ Analysis
 - Why?
 - What data did I use?
 - Structure of network
 - How I ran and compared
 - Transfer learning approach
- ◎ Results
 - Accuracy
 - Training Time
 - Other considerations
 - Ease of use
 - Documentation
 - job listings

Who am I?

Email: kevin.m.loew@gmail.com
Linkedin: www.linkedin.com/in/kmloew
GitHub: <https://github.com/kmloew>



Bachelor of Science
Embry-Riddle Aeronautical University
Research: Gravitational-Waves (LIGO)



Master of Science
Brandeis University
Research: High-Energy (CERN)



Doctor of Philosophy (Fall 2019)
Colorado State University
Research: Pattern Formation and nanostructures
(Bradley Theory Group)




This project can be found at: GitHub: https://github.com/KMLoew/DL_framework_comparison/tree/master

The data is from <https://www.kaggle.com/puneet6060/intel-image-classification>

Artificial Neural Network

Deep neural network



Artificial Neural Network

- ◎ Learn tasks with few (if any) specific rules
 - Like image classification
- ◎ Collection of artificial neurons and connections
- ◎ Multiple frameworks
 - TensorFlow
 - PyTorch
 - Sonnet
 - Gluon
 - Swift
 - etc...

What are TensorFlow and PyTorch

TensorFlow vs. PyTorch

- © Built on Theano
 - Google
 - © Static Graph
 - Keras is now part of TensorFlow and bypasses most of this
 - © More documentation!
 - © Lots of tutorials
- © Extended from Torch
 - Facebook
 - © Dynamic graph
 - © More “pythonic”
 - © Different tools
 - © Becoming more popular on Udemy and online courses

Which one?

- © No clear answer...
- © Different tools...



Bake Off!

- ◎ Solve the same problem in both
- ◎ Follow tutorials and examples to build and train simple network
- ◎ After 15 epochs what is the accuracy?
- ◎ Transfer learning?
 - Resnet50 pretrained on ImageNet data
 - 10 epochs→unfreeze→20 epochs

Intel Image Classification

<https://www.kaggle.com/puneet6060/intel-image-classification>



Building



Forest



Glacier



Mountain

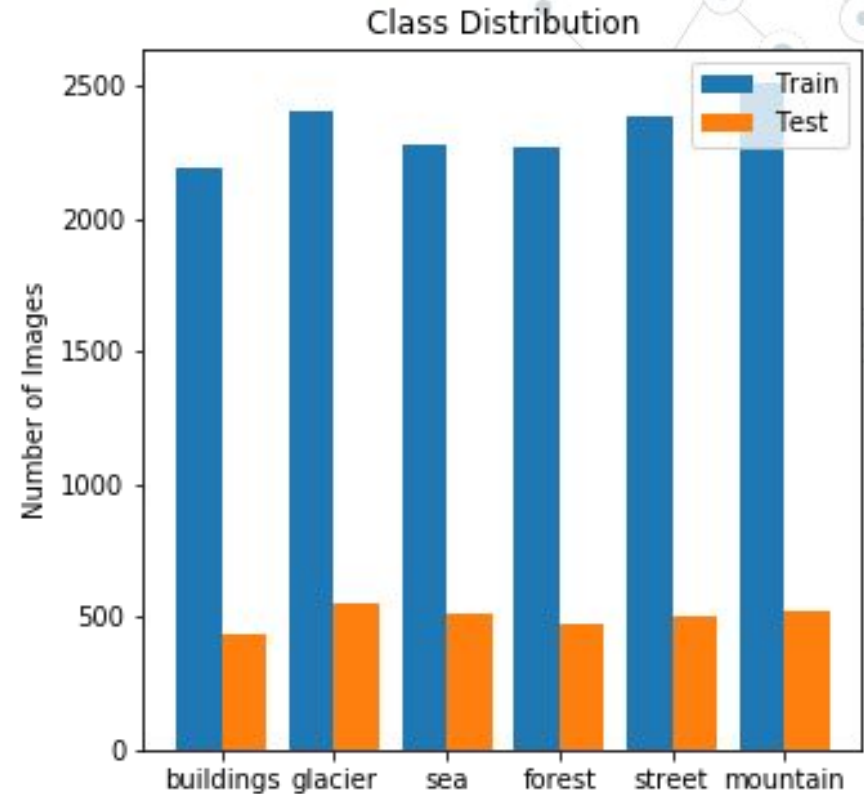
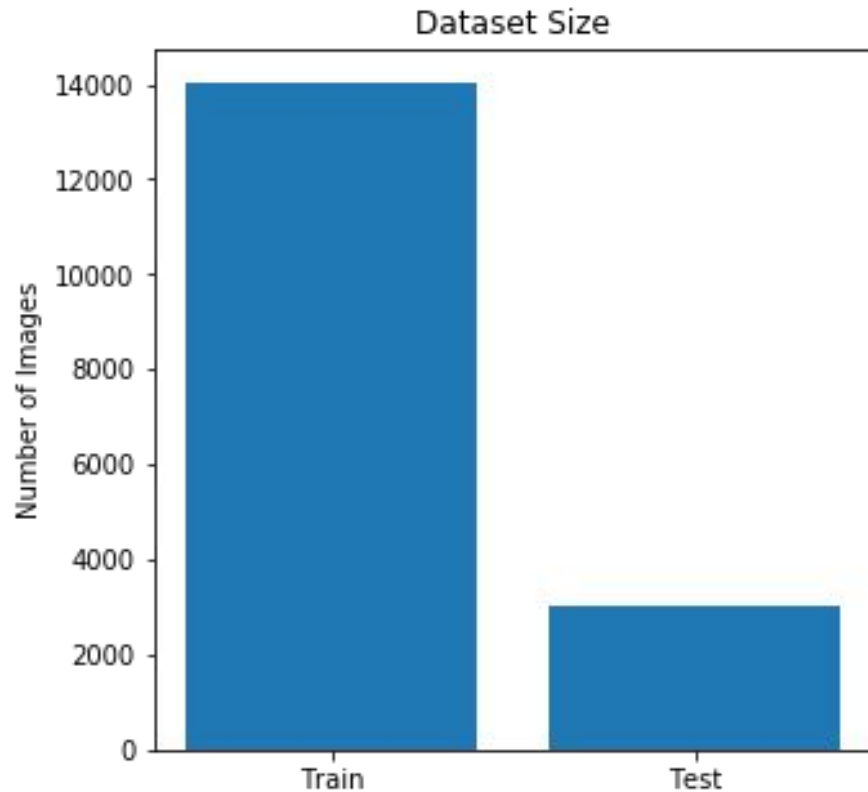


Sea



Street

Intel Image Classification



- ◎ 6 Classes
- ◎ Fairly evenly distributed
- ◎ Not all images are the same size

Hardware

© HP Omen

- Ubuntu 18.04
- Python: 3.6.8
- TensorFlow: 1.13.1
- Torch: 1.1.0
- Run locally
- No other applications running

Processor: Intel i7 8750H

Graphics: Nvidia GeForce GTX 1070


RAM: 32GB of RAM

Storage: 512GB SSD + 1TB Hard Drive




Simple Network

- ⊙ Convolutional Layer
- ⊙ Pooling Layer
- ⊙ Convolutional Layer
- ⊙ Pooling Layer
- ⊙ Flattening Layer
- ⊙ Dense (Fully connected) Layer
- ⊙ Dense (Fully connected) Layer
- ⊙ Dropout: 50% dropped connections
- ⊙ Dense (Fully connected) Layer




Training

- ◎ Accuracy as metric
- ◎ Sparse Categorical Cross Entropy as loss function
- ◎ Adam optimizer with learning rate of 0.001
 - fastai has variable learning rate
 - For the fine tuning of the TF resnet, the learning rate was 0.0001
- ◎ Batch size of 64



Data Augmentation

- ◎ Resize images to be 150x150
- ◎ Normalize pixel range to [0,1]
- ◎ Randomly rotate up to 25°
- ◎ Randomly flip images horizontally and vertically



TensorFlow+Keras Layers

```
1 model=keras.Sequential()
2
3 model.add(Conv2D(32, kernel_size=(3,3), strides=2, activation='relu', input_shape=(150,150,3)))
4 model.add(MaxPooling2D(pool_size=(2,2)))
5
6 model.add(Conv2D(64, kernel_size=(3,3), strides=2, activation='relu'))
7 model.add(MaxPooling2D(pool_size=(2,2)))
8
9 model.add(Flatten())
10 model.add(Dense(128,activation='relu'))
11 model.add(Dense(64,activation='relu'))
12 model.add(Dropout(rate=0.5))
13 model.add(Dense(len(classes), activation='softmax'))
```



- ◎ Very easy to build
- ◎ Lots of examples and documentation

PyTorch+nn Layers

```
1 class myModel(nn.Module):
2     def __init__(self):
3         super(myModel,self).__init__()
4
5         self.cnn1=nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=2)
6         self.relu=nn.ReLU()
7         self.maxpool1=nn.MaxPool2d(kernel_size=2)
8
9         self.cnn2=nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=2)
10        self.maxpool2=nn.MaxPool2d(kernel_size=2)
11
12        self.fc1=nn.Linear(in_features=64*9*9, out_features=128)
13        self.fc2=nn.Linear(in_features=128, out_features=64)
14        self.dropout=nn.Dropout(p=0.5)
15        self.fc3=nn.Linear(in_features=64, out_features=len(classes))
16
17    def forward(self,x):
18        out=self.cnn1(x)
19        out=self.relu(out)
20        out=self.maxpool1(out)
21        out=self.cnn2(out)
22        out=self.relu(out)
23        out=self.maxpool2(out)
24        out=out.view(64,64*9*9)
25        out=self.fc1(out)
26        out=self.relu(out)
27        out=self.fc2(out)
28        out=self.relu(out)
29        out=self.dropout(out)
30        out=self.fc3(out)
31        return out
```

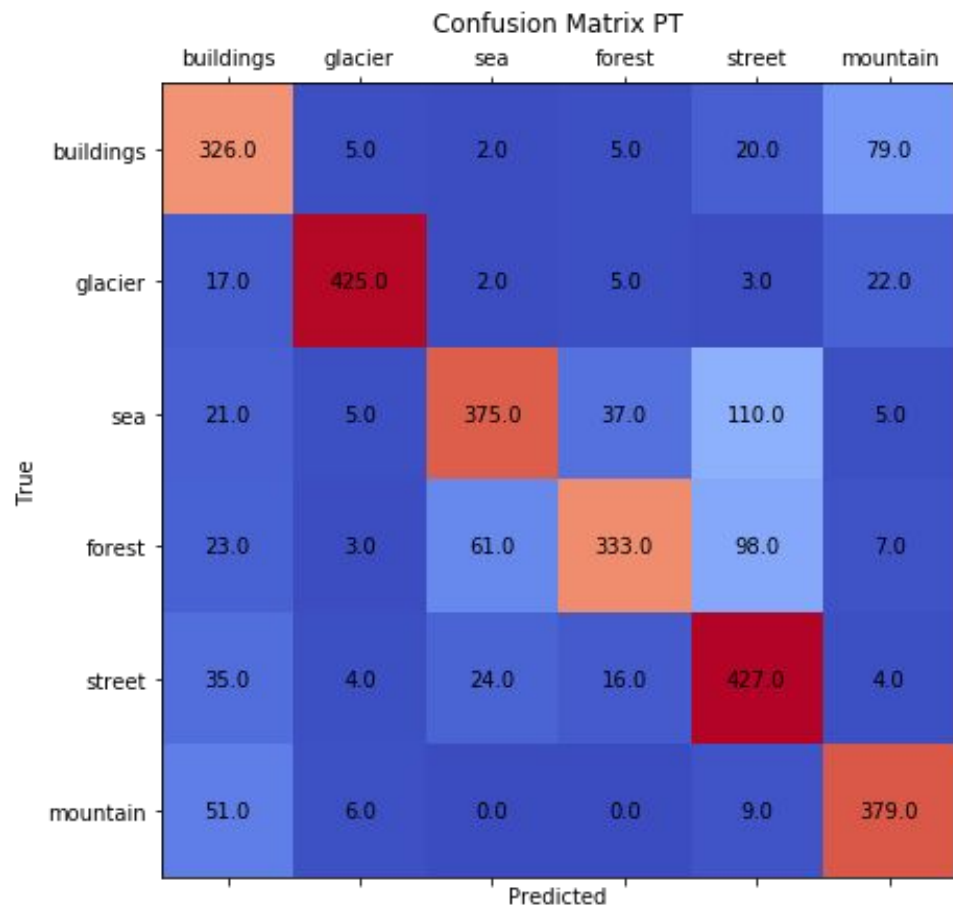
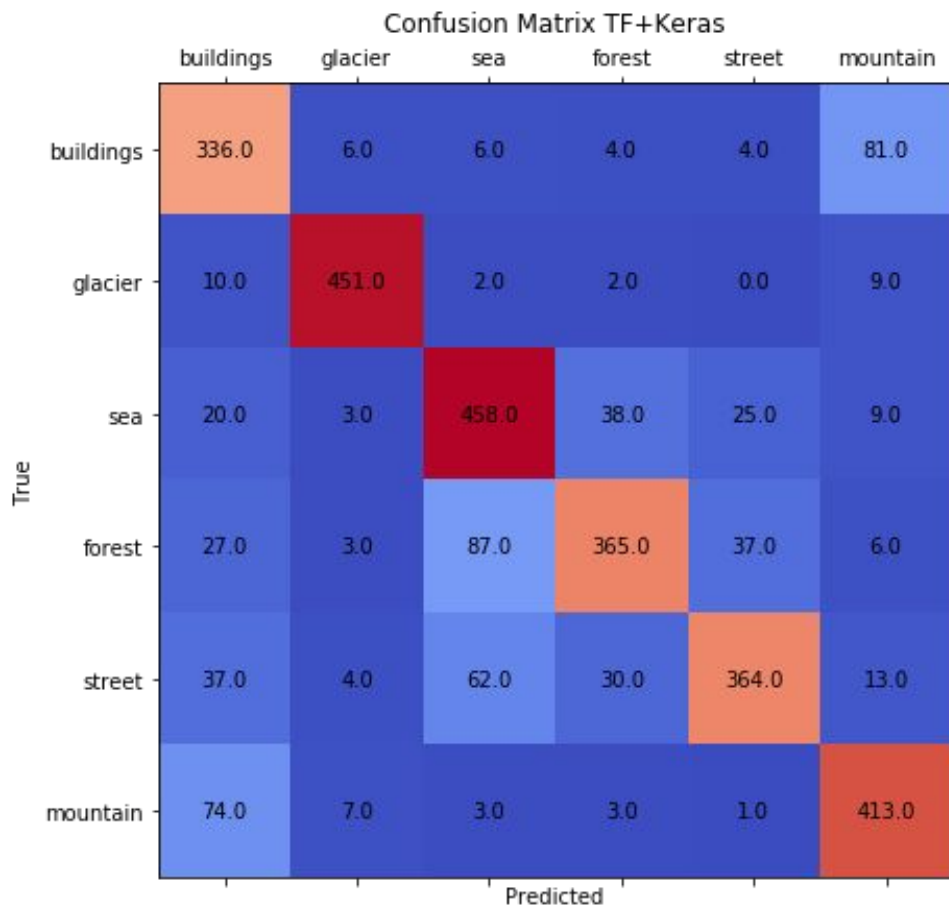
© Need to define more, harder to figure out...

Results: Training plots




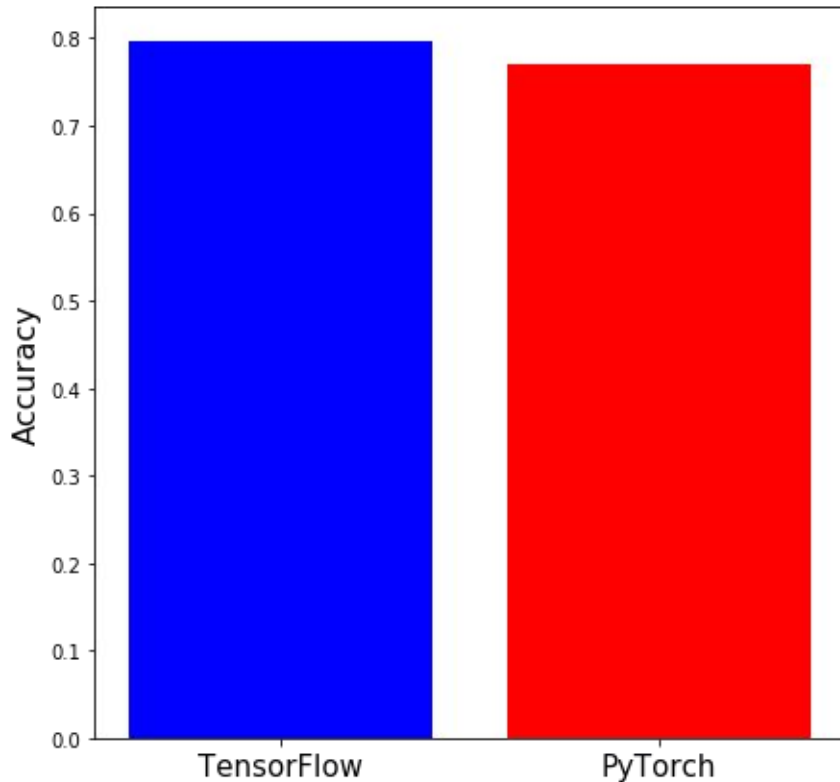
- © Similar accuracy curves
- © TensorFlow+Keras achieve higher accuracy

Results: Confusion Matrix



© Tensorflow with Keras layers does a bit better

Results: Accuracy and Training Time



- ◎ TensorFlow was a bit more accurate
 - 80% vs 77%
- ◎ PyTorch trained slightly faster
- ◎ Keras Layers were easier to work with

Transfer Learning

- ◎ Often instead of building a new network, an old one will be retrained
- ◎ This speeds up training because it already knows how to extract information, just needs to learn what you want
- ◎ “Cut off” top layers and place new ones
- ◎ I chose resnet50 as my base model
- ◎ Pretrained on Imagenet data (~14 million images)

TensorFlow

```
1 base_model=ResNet50(input_shape=(150,150,3),include_top=False)
```

```
1 model=keras.Sequential([  
2     base_model,  
3     GlobalAveragePooling2D(),  
4     Dense(len(classes),activation='softmax')  
5 ])   
6   
7 model.summary()
```


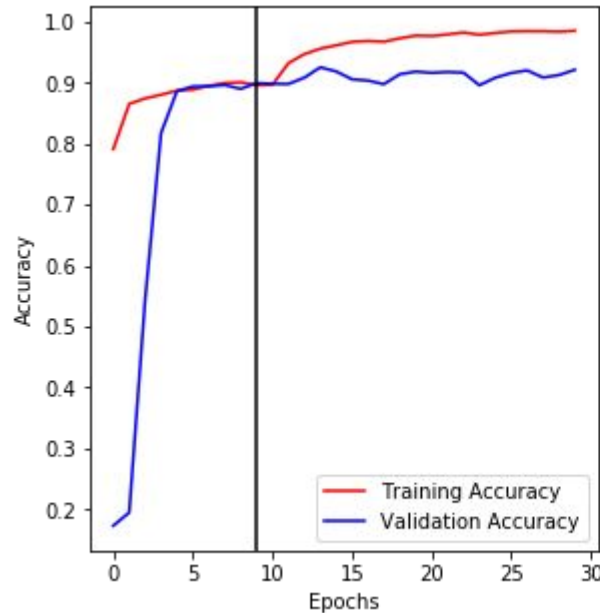
- ◎ Need to declare what layers are trainable
 - `base_model.trainable=False`
- ◎ Not too bad

Pytorch+fastai

```
1 model=cnn_learner(data,models.resnet50,metrics=accuracy,model_dir="/tmp/model/")
```

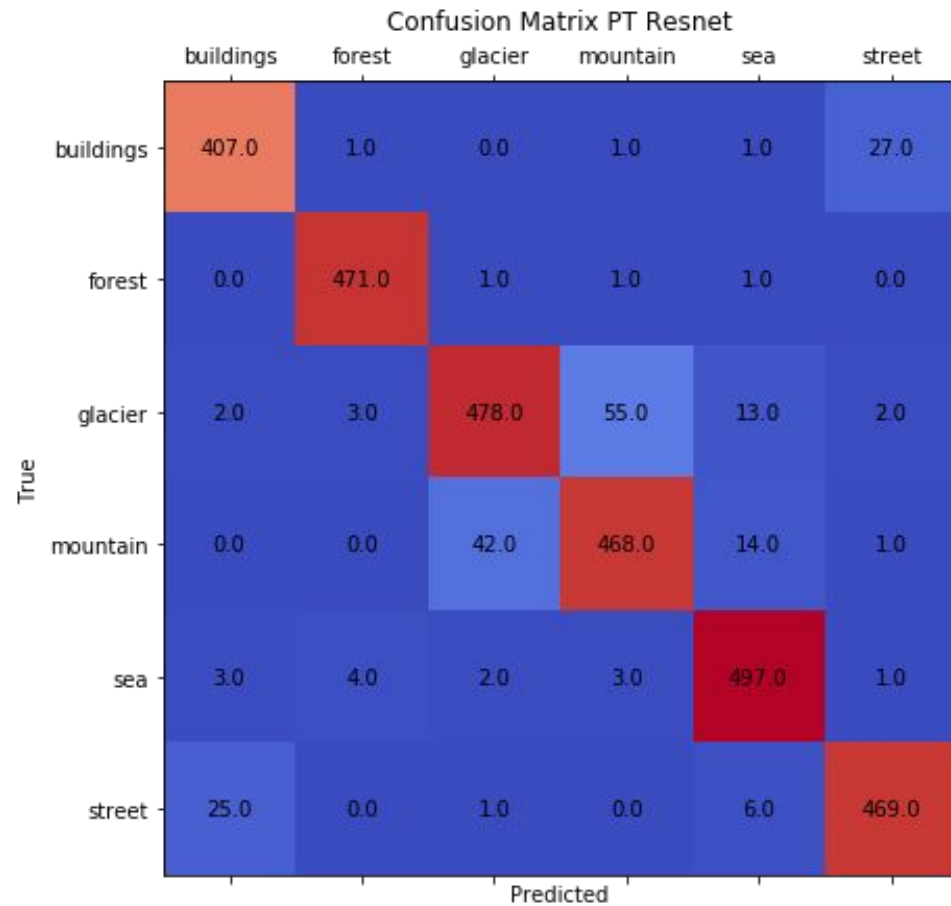
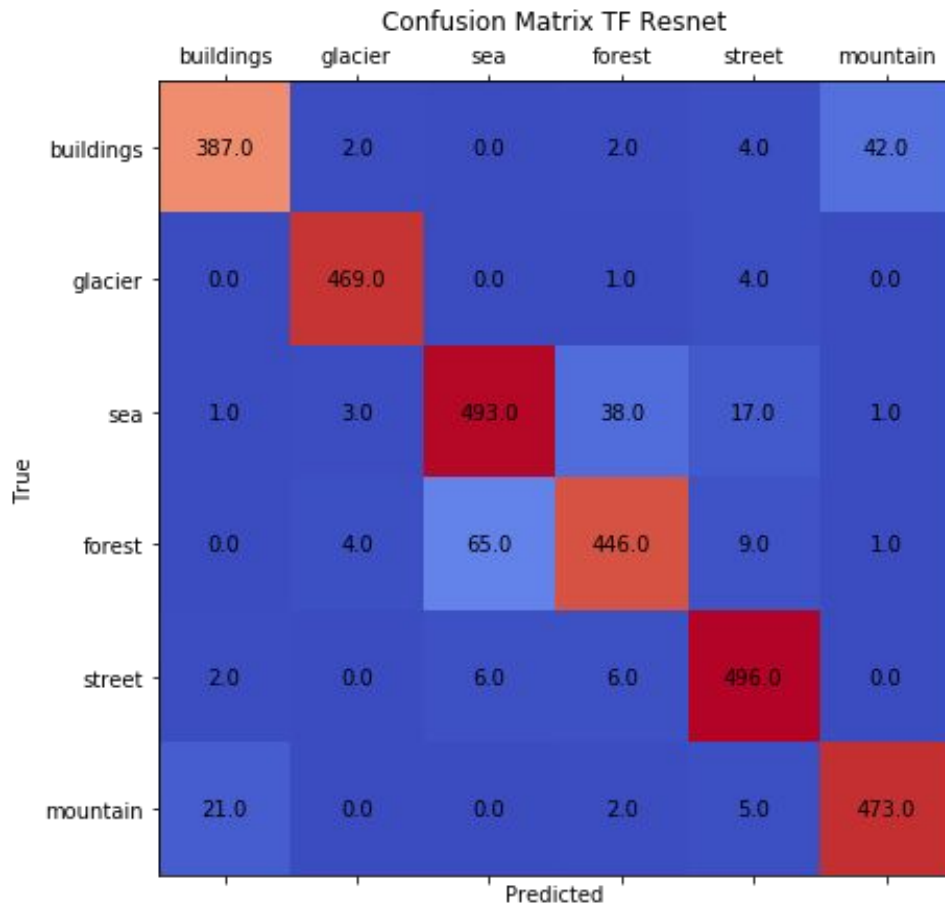
© So easy!

Results: Training plots



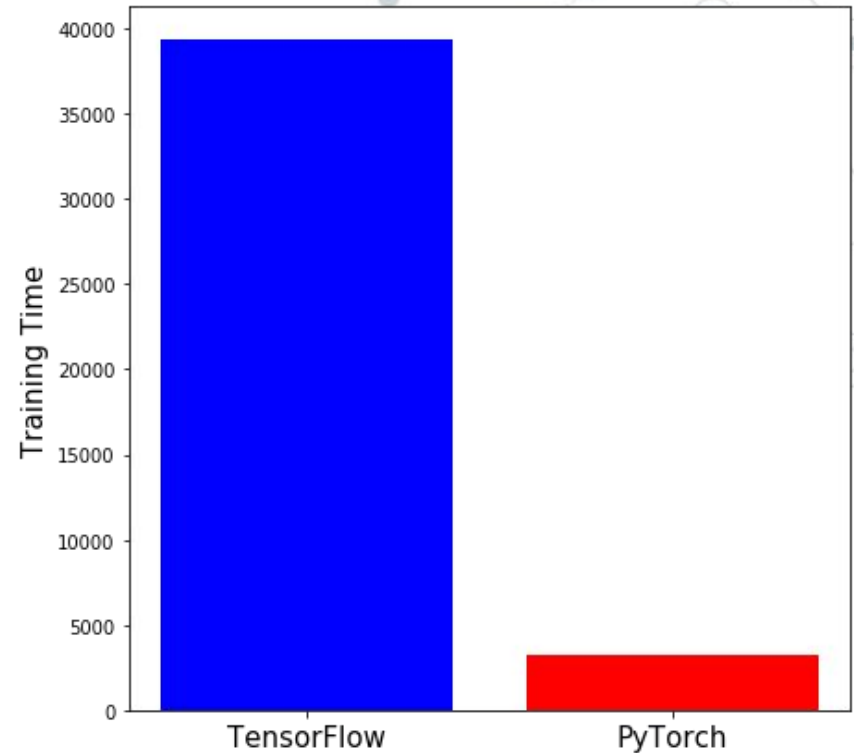
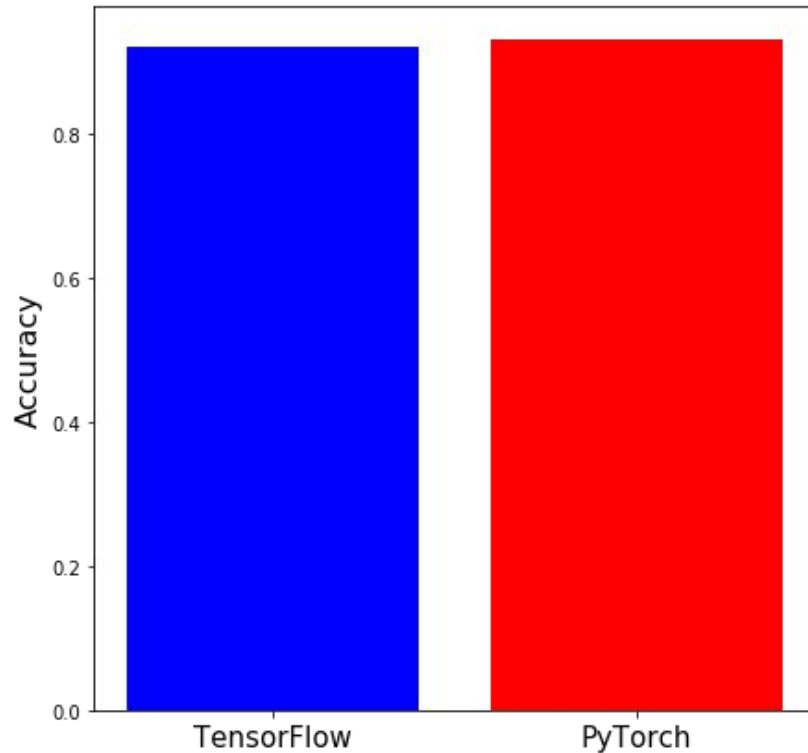
- © Both get higher accuracy than the simple model
- © PyTorch+fastai has more fluctuation for fine tuning

Results: Confusion Matrix



© Very similar

Results: Accuracy and Training Time



- ◎ Virtually Identical accuracy
 - 92% vs 93%
- ◎ The difference is in training time...

Final Results

Simple

- ◎ TensorFlow+Keras Layers were easier to build
- ◎ TensorFlow+Keras got higher final accuracy
- ◎ PyTorch Trained Faster
- ◎ Data handling was easier in PyTorch


PreTrained

- ◎ PyTorch+fastai was significantly faster
- ◎ PyTorch+fastai was easier to work with, once I figured it out

Other thoughts?

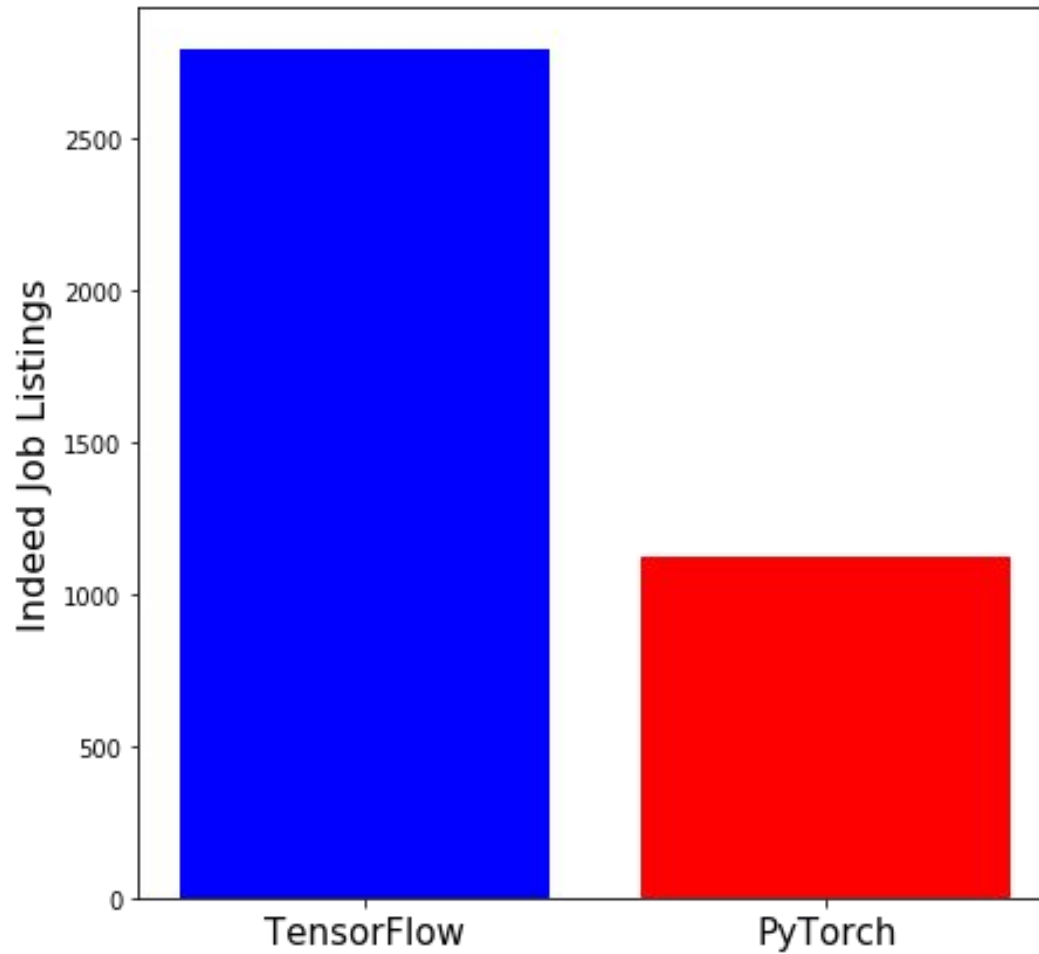
- © Which one is used in industry?
- © Amount of documentation?
- © Which is easier to learn?

Google Trends



TensorFlow
PyTorch

Indeed Job Postings



© Same trend as google

Conclusion

- ◎ Use the tool that makes sense for you
- ◎ TensorFlow is more popular and has more documentation
 - I would recommend this since it seems to be more common
- ◎ PyTorch is easy to pick up and experiment with
 - Good for learning concepts
- ◎ Fastai does amazing for transfer learning!